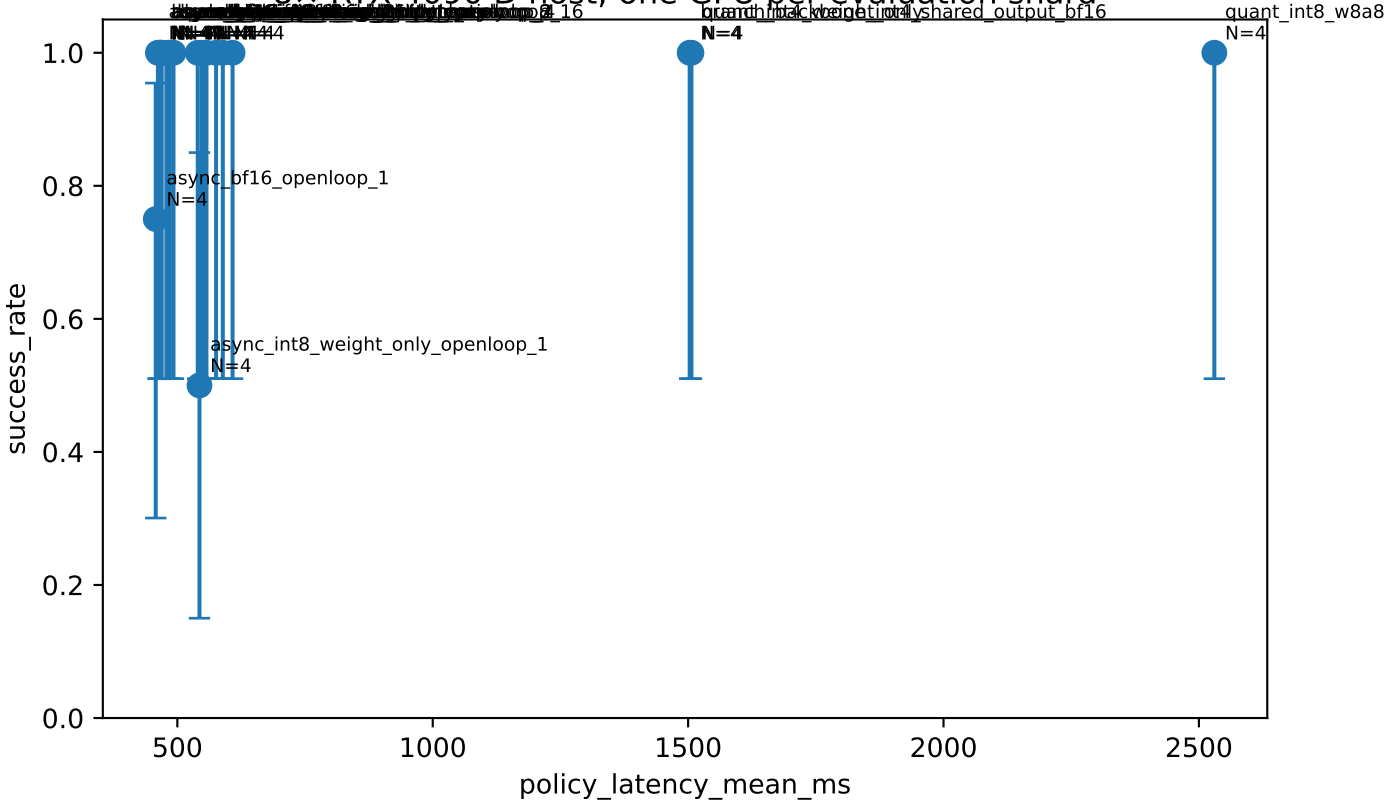

01 success vs policy latency
8× RTX 4090 D host; one GPU per evaluation shard
success_rate
policy_latency_mean_ms
0.0
0.2
0.4
0.6
0.8
1.0
500
1000
1500
2000
2500
async_bf16_openloop_1
N=4
async_int8_weight_only_openloop_1
N=4
quant_int8_w8a8
N=4