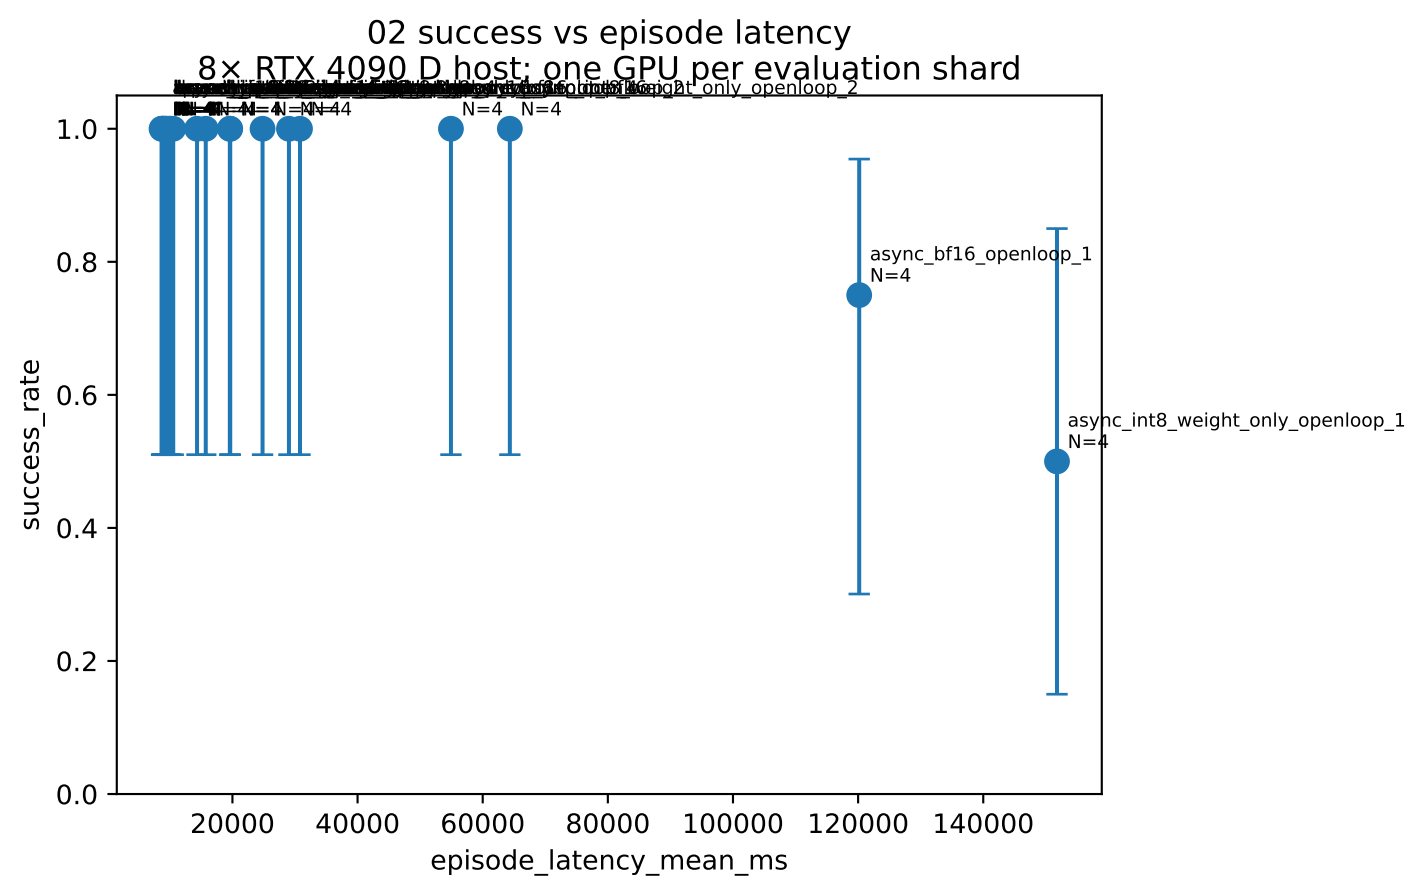
02 success vs episode latency
8× RTX 4090 D host; one GPU per evaluation shard
success_rate
episode_latency_mean_ms
0.0
0.2
0.4
0.6
0.8
1.0
20000
40000
60000
80000
100000
120000
140000
async_bf16_openloop_1
N=4
async_int8_weight_only_openloop_1
N=4
N=4
N=4