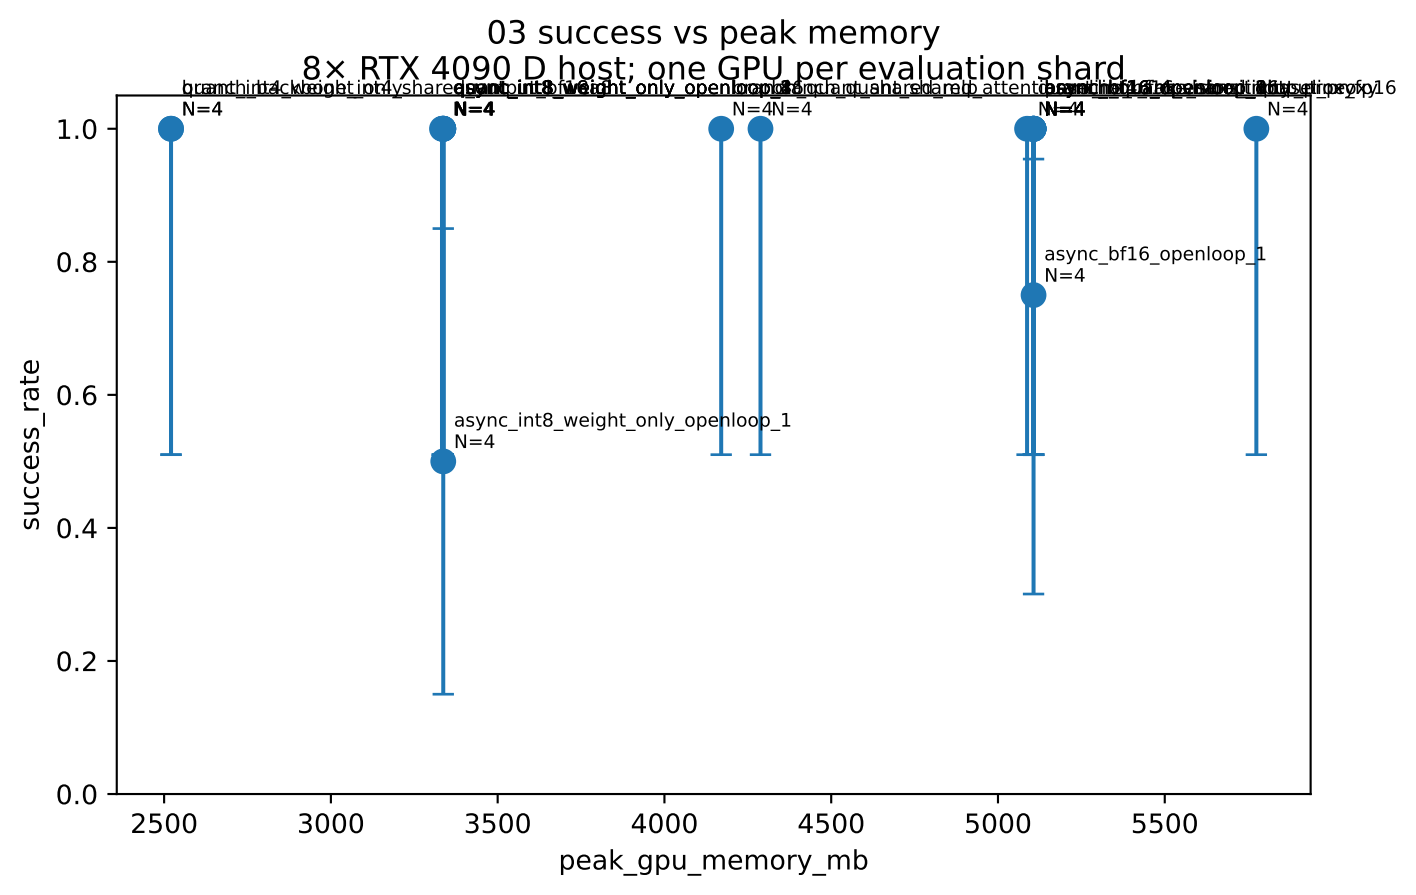

03 success vs peak memory
8× RTX 4090 D host; one GPU per evaluation shard
async_int8_weight_only_openloop_1
N=4
async_bf16_openloop_1
N=4
success_rate
peak_gpu_memory_mb
0.0
0.2
0.4
0.6
0.8
1.0
2500
3000
3500
4000
4500
5000
5500