

Precision versus policy denoising latency
actual fixed-observation microbenchmark, N=20-30 calls/mode, RTX 4090 D

