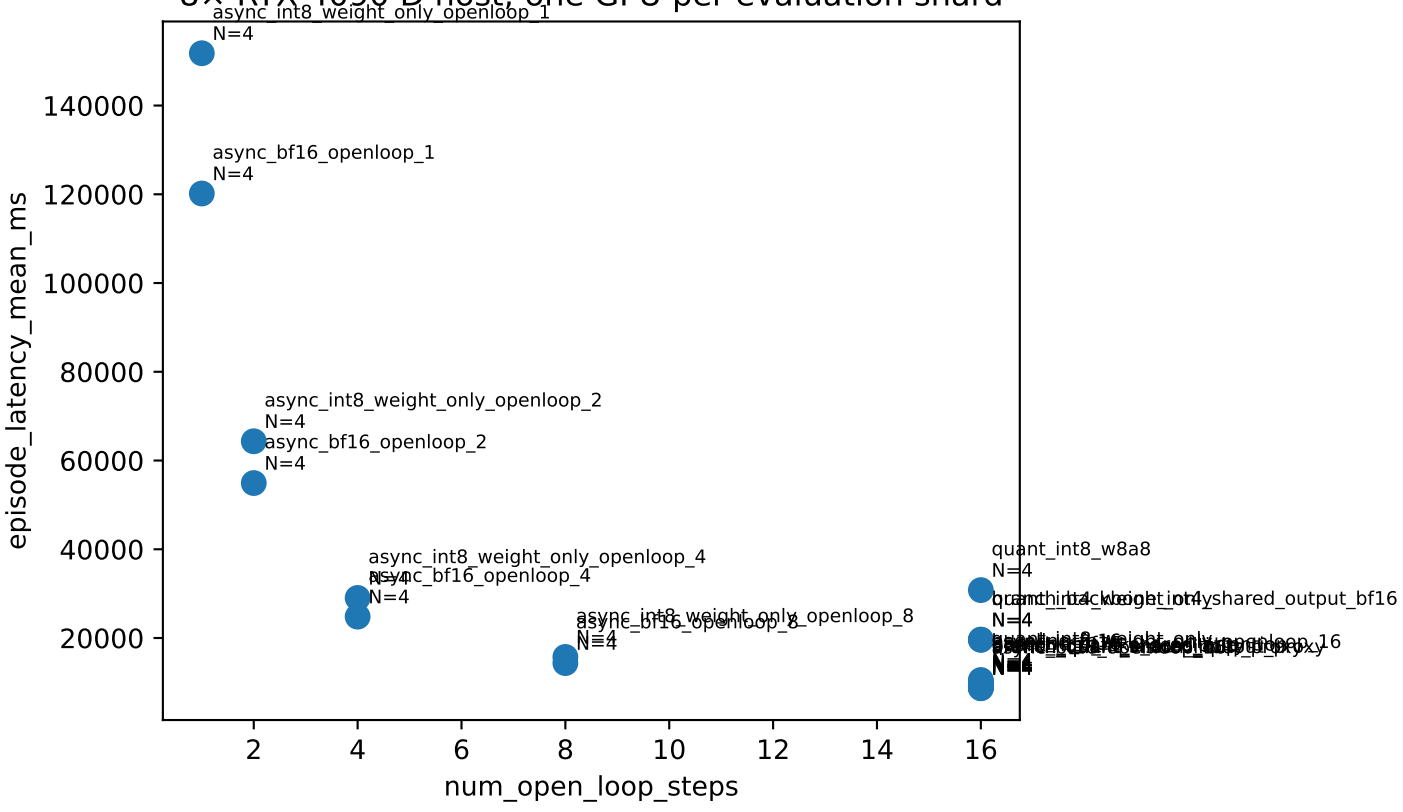
07 openloop vs episode latency
8× RTX 4090 D host; one GPU per evaluation shard
episode_latency_mean_ms
num_open_loop_steps
140000
120000
100000
80000
60000
40000
20000
2
4
6
8
10
12
14
16
async_int8_weight_only_openloop_1
N=4
async_bf16_openloop_1
N=4
async_int8_weight_only_openloop_2
N=4
async_bf16_openloop_2
N=4
async_int8_weight_only_openloop_4
N=4
async_bf16_openloop_4
N=4
async_int8_weight_only_openloop_8
N=4
async_bf16_openloop_8
N=4
quant_int8_w8a8
N=4
N=4